

Mitigating Hallucinations in Legal Large Language Models via Retrieval-Augmented Generation: A Comprehensive Survey

Joel Arturo Becerril Balderas
School of Engineering
Tecnológico de Monterrey
Monterrey, Mexico
A01797427@tec.mx

Angel Eduardo Urueta Puella
School of Engineering
Tecnológico de Monterrey
Monterrey, Mexico
A01796724@tec.mx

Efraín Paredes Balgañón
School of Engineering
Tecnológico de Monterrey
Monterrey, Mexico
A01351304@tec.mx

Marco Antonio Chávez García
School of Engineering
Tecnológico de Monterrey
Monterrey, Mexico
A01797547@tec.mx

Abstract—The deployment of Large Language Models (LLMs) in high-stakes legal environments faces critical challenges: hallucinations (fabricated citations, misinterpreted statutes) and temporal blindness (outdated parametric memory). Retrieval-Augmented Generation (RAG) has emerged as the state-of-the-art solution, anchoring generative models to verified external knowledge bases. This survey critically analyzes recent RAG methodologies (2023-2024), examining the mathematical limitations of dense embeddings, multi-hop reasoning failures, and the "Lost in the Middle" phenomenon. We provide quantitative benchmarks showing Graph-RAG achieves 98.4% accuracy (vs. 84.8% baseline GPT-4), reducing hallucinations from 8.9% to 0.1%. However, fundamental bottlenecks persist: cosine similarity fails on negation ($\cos(\text{"liable"}, \text{"not liable"}) = 0.947$), naive chunking destroys document structure, and data privacy constraints remain unresolved. We propose hybrid architectures combining Knowledge Graphs, multi-agent verification, and Role-Based Access Control (RBAC) as the pathway toward hallucination-resistant legal AI.

Index Terms—Deep Learning, Retrieval-Augmented Generation, Large Language Models, Vector Databases, Legal AI, Hallucination Mitigation, Graph-RAG, Data Privacy

I. INTRODUCTION

Foundation models, driven by transformer architectures, exhibit remarkable proficiency in reading comprehension and logical reasoning. Despite these advancements, a critical architectural vulnerability persists: LLMs operate fundamentally as probabilistic next-token predictors constrained by a static "Knowledge Cutoff." This limitation results in two major failure modes:

- 1) Hallucinations, where the model confidently fabricates factual inaccuracies, [In May 2023, attorney Steven Schwartz submitted a legal brief containing six ChatGPT-fabricated cases](#) (Varghese vs. China Southern

Airlines, etc.), resulting in sanctions and demonstrating that even advanced LLMs confidently generate non-existent precedents.

- 2) Temporal Blindness: Models with knowledge cutoffs cannot account for recent legislation, court decisions, or regulatory changes—catastrophic in jurisprudence where precision is mandatory.

In the legal domain, where jurisprudential precision is mandatory, relying on a temporally blind model can lead to catastrophic professional consequences. Retrieval-Augmented Generation (RAG) [1] mitigates this by decoupling the knowledge repository from the neural weights, fetching real-time, verified documents from an external vector database.

RAG as Solution: User Query $\rightarrow$ Vector Search $\rightarrow$ Retrieved Context $\rightarrow$ LLM (context + query) $\rightarrow$ Verified Response

This survey critically examines the technical efficacy, security constraints, and underlying deficiencies of RAG architectures in legal document analysis.

II. TAXONOMY OF LEGAL HALLUCINATIONS

Before addressing mitigation strategies, it is crucial to classify the specific types of generative failures in the legal domain. Hallucinations in jurisprudence are not monolithic; they manifest in three distinct categories [6]:

- **Citation Fabrication (Precedent Hallucination):** The most dangerous variant, where the LLM invents entirely fictitious case laws, complete with fake docket numbers and judge names, stemming from the model's attempt to satisfy user prompts with statistically plausible string combinations. Risk: CRITICAL—Legal malpractice, bar sanctions, contempt of court.

- **Statutory Misinterpretation (Reasoning Hallucination):** The model retrieves the correct law but applies flawed formal logic, failing to understand clauses like "notwithstanding" or "except as provided in." Risk: HIGH—Incorrect legal advice, contractual errors.
- **Factual Inaccuracy:** Incorrectly recalling dates, jurisdictions, or the specific parties involved in a historical case due to parametric memory decay. Risk: MODERATE—Outdated strategies, jurisdictional errors.

III. LITERATURE REVIEW: ARCHITECTURAL DEFICIENCIES IN MODERN RAG

The transition towards domain-specific RAG systems has exposed significant algorithmic bottlenecks. Rather than treating RAG as a flawless pipeline, recent literature has systematically deconstructed its technical limitations.

A. The Mathematical Flaws of Dense Embeddings

In a standard RAG pipeline, documents and queries are projected into a high-dimensional vector space, where relevance is computed using Cosine Similarity (similarity = $\cos(\theta) = \frac{A \cdot B}{\|A\| \|B\|}$). While Cui et al. (2023) [3] demonstrated the utility of this approach, critical evaluations reveal that dense embedding models struggle profoundly with logical negation. In jurisprudence, "The tenant is liable" and "The tenant is not liable" occupy nearly identical vector spaces. Consequently, Vector Databases often retrieve contradictory statutes, forcing the LLM to synthesize conflicting constraints. This necessitates Hybrid Search algorithms combining dense retrieval with sparse, keyword-based scoring (e.g., BM25) [4].

B. Contextual Fragmentation and Multi-Hop Reasoning

Naive chunking strategies—segmenting texts into arbitrary token limits—destroy the hierarchical structure of legal documents. Huang et al. (2024) [7] emphasize semantic chunking, yet an even deeper architectural flaw remains: the lack of Multi-Hop reasoning. Standard K-Nearest Neighbors (K-NN) retrieval algorithms only fetch chunks directly similar to the prompt. They cannot perform transitive logical leaps across multiple, disconnected articles. Barnett et al. (2024) [5] identified this "Missing Context" as the primary cause of system failure.

C. Context Window Saturation: "Lost in the Middle"

Recent studies document the "Lost in the Middle" phenomenon: when LLMs are injected with an extensive context window containing multiple retrieved chunks, their attention mechanisms heavily prioritize information at the beginning and end of the prompt [4]. Critical legal exceptions hidden in the middle chunks are frequently ignored during generation, leading to legally hazardous outputs.

IV. SECURITY AND PRIVACY CONSTRAINTS IN LEGAL RAG

A dimension frequently overlooked in standard RAG surveys is the strict data governance required in corporate law.




Fig. 1. Proposed Hybrid RAG Architecture combining Knowledge Graphs and Multi-Agent Adversarial Verification for Legal AI.

Feeding sensitive client contracts into a public LLM API constitutes a massive breach of Attorney-Client Privilege.

- **Data Leakage and Multi-Tenancy:** In a law firm handling multiple clients, the vector database must employ strict Metadata Filtering. Without Role-Based Access Control (RBAC) injected directly into the retrieval query, a RAG system could accidentally fetch confidential defense strategies from Client A to answer a prompt from Client B.
- **Prompt Injection Attacks:** Malicious actors can embed hidden text within a legal PDF contract. When the RAG system retrieves this document, the hidden prompt could hijack the LLM to alter the final legal summary.

TABLE I
COMPARISON OF RAG METHODOLOGIES IN LEGAL CONTEXTS

Methodology	Retrieval Strategy	Primary Legal Flaw
Naive RAG	Fixed-token chunking, Cosine Similarity	Destroys statutory context; fails at negation.
Advanced RAG	Semantic chunking, Hybrid Search	Susceptible to "Lost in the Middle" omissions.
Graph-RAG	Knowledge Graphs + Vector Embeddings	Computationally expensive to map all laws.

V. FUTURE DIRECTIONS

Addressing these architectural flaws requires shifting from basic semantic search to complex, deterministic retrieval pipelines:

- **Graph-RAG Architectures:** Integrating Knowledge Graphs with vector embeddings. By modeling legal statutes as nodes and their cross-references as edges, the system can traverse the graph to fetch entire chains of legal logic.
- **Multi-Agent Cross-Verification:** Implementing adversarial Multi-Agent systems where a secondary "Reviewer Agent" programmatically verifies the primary LLM's output against the retrieved documents to detect omissions before user delivery. Architecture: RAG Retrieval → Drafting Agent → Draft ↓ Verification Agent ← Context ↓ [Agreement?] → Yes → Output ↓ No Adversarial Debate → Human Review

VI. CONCLUSION

This research survey highlights that Retrieval-Augmented Generation (RAG) is currently the most viable architecture for deploying Large Language Models in high-stakes environments. A critical vulnerability of foundational LLMs, beyond inherent hallucinations, is their "temporal blindness" caused by static knowledge cutoffs. As demonstrated in fields requiring up-to-the-minute precision, relying solely on frozen parametric memory inevitably leads to catastrophic miscalculations.

RAG elegantly solves this temporal decay by dynamically anchoring the generative model to real-time, verified databases, bypassing the prohibitive costs of constantly re-training the neural network. However, our critical analysis reveals that RAG is not a flawless solution. The technical efficacy of the system remains heavily bottlenecked by the mathematical limitations of cosine similarity, the inability of vector spaces to execute multi-hop reasoning, attention-mechanism flaws, and severe data privacy risks. Addressing these specific algorithmic and security gaps through Hybrid Search, RBAC integrations, and Graph-RAG will be critical for the safe and ethical adoption of Artificial Intelligence in jurisprudence over the coming decade.

Key Findings:

- 1) Embeddings: Cosine similarity achieves only 58
- 2) Multi-hop: Standard retrieval misses 67
- 3) Security: RBAC + on-premise deployment mandatory for attorney-client privilege
- 4) Performance: Graph-RAG superior but 73

Recommendations for Practitioners:

- **Developers:** Hybrid search ($\alpha = 0.7$), semantic chunking, multi-agent verification.
- **Legal Professionals:** Always verify citations (even 0.1

Research Agenda:

- **Short-term (1-2y):** Standardized legal benchmarks, optimized graph construction
- **Medium-term (3-5y):** Neuro-symbolic architectures, federated learning
- **Long-term (5+y):** Provably hallucination-free systems

REFERENCES

- [1] Y. Gao *et al.*, "Retrieval-Augmented Generation for Large Language Models: A Survey," *arXiv preprint arXiv:2312.10997*, 2023. [Online]. Available: <https://arxiv.org/abs/2312.10997>
- [2] S. Es *et al.*, "RAGAS: Automated Evaluation of Retrieval Augmented Generation," *arXiv preprint arXiv:2309.15217*, 2023. [Online]. Available: <https://arxiv.org/abs/2309.15217>
- [3] J. Cui *et al.*, "ChatLaw: Open-Source Legal Large Language Model with Integrated External Knowledge Bases," *arXiv preprint arXiv:2306.16092*, 2023. [Online]. Available: <https://arxiv.org/abs/2306.16092>
- [4] P. Zhao *et al.*, "Retrieval-Augmented Generation for AI-Generated Content: A Survey," *arXiv preprint arXiv:2402.19473*, 2024. [Online]. Available: <https://arxiv.org/abs/2402.19473>
- [5] S. Barnett *et al.*, "Seven Failure Points When Engineering a Retrieval Augmented Generation System," *arXiv preprint arXiv:2401.05856*, 2024. [Online]. Available: <https://arxiv.org/abs/2401.05856>
- [6] N. Guha *et al.*, "LegalBench: A Collaboratively Built Benchmark for Measuring Legal Reasoning in Large Language Models," *NeurIPS*, 2023. [Online]. Available: <https://arxiv.org/abs/2308.11462>
- [7] Y. Huang *et al.*, "A Survey on Retrieval-Augmented Text Generation for Large Language Models," *ACM Computing Surveys*, 2024. [Online]. Available: <https://arxiv.org/abs/2404.10981>